Google Just Turned Chrome Into a Programmable AI Assistant. Here's How to Use It.

Kristopher Dunham

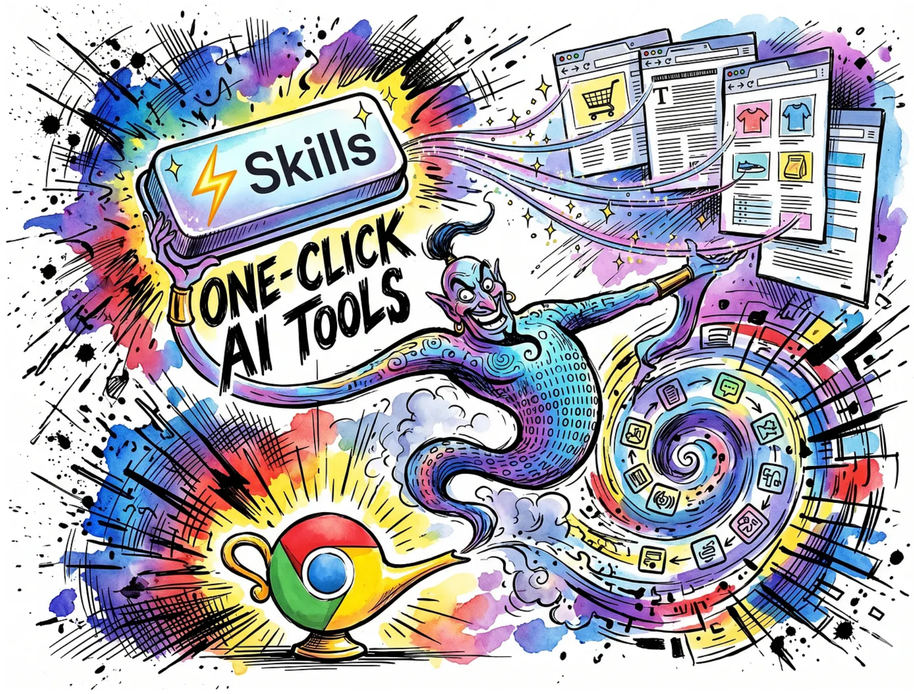

Every power user of AI has the same guilty routine: a text file somewhere full of prompts they copy and paste over and over again.

Maybe it's a prompt that breaks down ingredient lists on recipe pages. Maybe it compares product specs across tabs. Maybe it summarizes dense articles into three bullet points. Whatever it is, the workflow looks the same every time: open the AI sidebar, paste the prompt, wait, get results, close, repeat tomorrow. And the day after that.

Google just killed that entire loop. On April 14, 2026, Chrome shipped a feature called Skills that lets you save any Gemini prompt as a reusable, one-click tool that runs against whatever webpage you're looking at. Or across up to ten tabs at once.

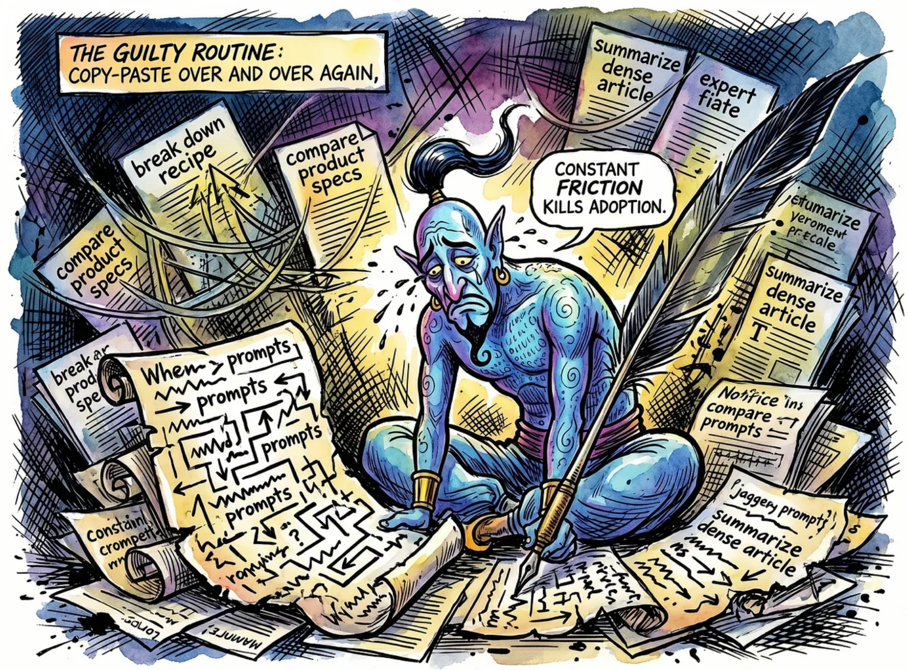

That last part is the real story. But we'll get there.

What Skills Actually Are (and Aren't)

Strip away the marketing language and Skills are surprisingly simple: they're saved prompts with a name and an emoji attached. You write a prompt in Gemini's Chrome side panel, get a good result, and hit "Save as Skill." Next time you need it, type `/` in the Gemini chat box, pick your

Skill from the dropdown, and it runs instantly against the current page.

That's it. No code. No extensions. No API keys.

But "simple" and "not useful" are very different things. Think about what this actually changes. Before Skills, using Gemini consistently meant remembering exactly how you phrased a prompt three weeks ago, or hunting through chat history to find the one that worked. The friction was small but constant, and constant friction kills adoption. You'd default back to doing things manually because the overhead of getting the AI to behave the right way, again, just wasn't worth it.

Skills remove that friction entirely. You optimize a prompt once, save it, and it becomes a permanent tool in your browser. It syncs across every desktop where you're signed into Chrome. You go from "I need to figure out how to ask this" to "I press one button."

The Multi-Tab Trick Nobody's Talking About

Here's where it gets genuinely interesting.

When you trigger a Skill, you can select additional open tabs to include in the context. Up to ten of them. The Skill runs against all of those pages simultaneously, pulling information from each one and synthesizing the results in a single output.

Think about what that means in practice. You're shopping for a laptop and you've got five retailer tabs open. Instead of manually comparing specs by toggling between windows, you save a "Spec Comparison" Skill once, select all five tabs, and get a unified comparison table in seconds. Each page might format its specs differently, use different terminology, bury the information in different places. The AI handles all of that normalization for you.

Or you're doing research. You have eight articles open on the same topic. A "Research Synthesis" Skill scans all of them, identifies the key arguments, notes where sources agree or disagree, and produces a summary that would have taken you an hour to compile manually.

This is the capability that transforms Skills from "nice convenience" to "genuinely different way of working." Your browser tabs stop being passive windows and start functioning as a live data corpus that the AI can query on demand.

How to Build Your First Skill (5 Minutes, Zero Code)

Let's make this concrete. Here's the exact process:

Step 1: Navigate to a page you work with regularly. A recipe blog, a product listing, a news article, whatever represents a task you repeat often.

Step 2: Open Gemini in Chrome. The fastest way is the keyboard shortcut: `Ctrl + G` on Windows/Linux or `Option + G` on Mac. You can also click the Gemini icon in Chrome's toolbar.

Step 3: Write a detailed prompt. Don't be vague. The more specific you are about what you want extracted, how you want it formatted, and what you want ignored, the better the Skill will perform across different pages.

Here's an example that works well for product research:

> "Analyze this product page. Extract the price, key technical specifications, return policy, and any mentioned warranty details. Format the output as a clean table. Flag anything that seems like marketing language rather than a verified specification."

Step 4: Review the output. If it's not quite right, refine the prompt through conversation. Tell Gemini what to fix. The system is designed for this iterative loop. Get the output exactly how you want it before saving.

Step 5: Click "Save as Skill." Give it a short, memorable name like "Product Breakdown" and assign an emoji. That's your new tool.

Step 6: Test it. Navigate to a completely different product page, type `/`, select your Skill, and see how it performs on unfamiliar content. If it needs adjustment, click the compass icon in the Gemini panel to open the management dashboard and edit the underlying prompt.

The whole process takes about five minutes. And you only do it once.

The Pre-Built Library: Start Here If You're Not Sure What to Build

Google ships Chrome with a library of ready-made Skills at `chrome://skills/browse`. These are worth browsing even if you plan to build your own, because they demonstrate what good prompt templates look like and give you a starting point to customize.

A few that stand out from the initial collection:

Protein Maximizer pulls nutritional data from recipe pages and recalculates for higher-protein variations. If you've ever tried to adapt an online recipe to hit specific macro targets, you know how tedious that math gets.

Spec Comparison is built specifically for multi-tab use. Open a few product pages, trigger this Skill, and get a side-by-side table without copying a single data point manually.

Ingredient Decoder breaks down unfamiliar chemical names on product or skincare pages into plain language with safety information. This one is genuinely useful for anyone who's ever stared at a label and wondered what they're actually buying.

Accessibility Testing Superpower runs a WCAG compliance check against whatever page you're viewing and outputs a remediation checklist. Developers and designers, take note. This one alone could save hours of manual auditing per week.

The key detail about pre-built Skills: once you add them to your collection, they're fully editable. You can take Google's "Ingredient Decoder" and modify the prompt to focus specifically on allergens, or adjust the output format to match how you prefer to read information. The library is a starting point, not a locked-down product.

The Safety Model: Why It Asks Before Sending That Email

Here's a question that probably crossed your mind: if Skills can interact with Gmail and Google Calendar, what stops a poorly written prompt from accidentally sending emails or booking meetings?

Short answer: confirmation gates.

When a Skill tries to perform what Google calls a "high-consequence action," the system pauses and shows you exactly what it's about to do. You see the drafted email, the proposed calendar event, the specific action it wants to take. Nothing happens until you explicitly click to confirm.

This matters more than it sounds. The whole point of Skills is speed and automation. You're running saved prompts quickly, probably with less scrutiny than a manually typed request. Google clearly thought about the failure mode where someone triggers a Skill without fully thinking through the consequences, and built in the right guardrails.

One detail worth knowing: some actions cascade in ways that can't be undone. If a Skill modifies a calendar event and that event has guests, the notification emails to those guests fire immediately. The calendar change itself might be reversible, but the notifications aren't. The confirmation gate is your chance to catch that before it happens.

Who Should Care About This

If you do any kind of research online, Skills change your workflow. Save prompts for summarizing articles, extracting key claims, comparing sources across tabs, or pulling structured data from unstructured pages. The multi-tab capability alone is worth the setup time.

If you shop online frequently, the comparison and analysis Skills save real time and probably real money. Having an AI break down specifications, flag marketing fluff, and normalize data across competing products means you make better decisions faster.

If you build websites or apps, the accessibility testing Skill and the ability to write custom analysis prompts that run against any page gives you a lightweight QA tool that lives right in your browser.

If you manage content or do SEO work, you can build Skills that audit title tags, check heading structures, or extract metadata across competitor pages in a single pass.

And if you're someone who's been curious about AI but never quite found a use case that stuck, Skills might be the thing that clicks. The friction of prompt writing was a real barrier for a lot of people. Removing it, and providing a library of proven templates, lowers the bar significantly.

The Bigger Picture: Your Browser Is Becoming an Agent

Zoom out for a second and look at the trajectory.

A year ago, Chrome's Gemini integration was basically a chatbot that could read the current page. Earlier this year, Google added multi-tab context and connections to Gmail and Calendar. Now Skills add reusability and workflow automation. And running in parallel is an experimental feature called Auto Browse that lets Gemini physically navigate websites, click buttons, and fill out forms on your behalf.

Connect those dots. A Skill that today reads a recipe and recalculates the macros could tomorrow navigate to a grocery delivery site, add the modified ingredients to a cart, and checkout for you. The confirmation gate would still fire before the purchase, but the workflow from "I found a recipe I like" to "groceries are on the way" becomes a single saved tool.

That's the trajectory Google is building toward. The browser isn't becoming a smarter search engine. It's becoming a programmable agent that takes action on your behalf, with your explicit approval at every consequential step.

Whether that excites or unnerves you probably depends on how much you trust the guardrails. But the direction is clear. And Skills are the foundation it's built on.

Getting Started Today

Here's what you need:

- Chrome desktop (Mac, Windows, or ChromeOS), updated to the latest version
- Signed into a Google Account
- Browser language set to English (US)
- No Incognito mode (Skills don't work there)

Google is still rolling this out in stages, so if you don't see Skills in your Gemini panel yet, give it a few days. These staged rollouts are standard for Chrome features, and availability tends to expand quickly once the initial wave lands.

Open the Gemini side panel with `Ctrl + G` (Windows) or `Option + G` (Mac). Browse the Skills library at `chrome://skills/browse` to grab a few pre-built options. Then write one custom Skill for a task you repeat at least weekly.

Start small. One Skill. Use it for a week. You'll build more once you see how much time the first one saves.

If you found this useful, I write regularly about AI tools, developer workflows, and the practical side of putting new technology to work. Follow along for more.

Artificial Intelligence · AI · AI Agent · Productivity · Google Gemini

Written by Kristopher Dunham

Follow

640 followers · 9 following

Published author, developer, creative who builds worlds. Fervorlife.com. Join the fun at CreativeAiDojo.com or CreativeAi.Ninja

Responses (1)

Pancake Agree

What are your thoughts?

See all responses